

# Efficient prime-field arithmetic in Rust


Simon Shine <simon@neptune.cash>

Slides available on:

<https://github.com/sshine/prime-field-benchmarks/>

# Who am I?

I'm Simon.

I work at Triton Software , and I make zero-knowledge cryptography for a privacy-focused blockchain called Neptune: <https://neptune.cash/>



# Who am I?

I'm Simon.

In my spare time, I study Chinese and practice time travel (mostly forwards, in kayaks).



# Disclaimer

I'm not a mathematician. Don't hesitate to add to what I say. Also, I didn't invent anything here. All of this is re-discovered. I try to give credit.

***tl;dr***

$| +$  and  $\times \bmod p$  faster using a combination of hacks, tricks, and knowledge of  $p$ .

# Field arithmetic

A **field** is a set  $S$  with  $+$  and  $\times$  where

- $+$  and  $\times$  associate, commute and distribute
- $\exists 0 \in S$  so for all  $a \in S$ ,  $0 + a = a$  and  $a + 0 = a$ .
- $\exists 1 \in S$  so for all  $a \in S$ ,  $1 \times a = a$  and  $a \times 1 = a$ .
- $\forall a \in S$ ,  $\exists b \in S$  so  $a + b = b + a = 0$
- $\forall a \in S$ ,  $a \neq 0$ ,  $\exists b \in S$  so  $a \times b = b \times a = 1$
- $0 \neq 1$

# Field arithmetic

It just means you can  $+$ ,  $-$ ,  $\times$ ,  $\div$  as you'd expect.

Reals and rationals are examples of infinite fields.

Most people call this "**math**".

# Finite-field arithmetic

An example of an (efficient) **finite field** is  
u64 with overflow (aka  $\text{GF}(2^{64})$ ).

Most programmers call this "**math**".



# Prime-field arithmetic

A **prime field**  $\mathbb{F}_p$  is a finite field that overflows at  $p$ .  
Or said with Rust code,

```
use std::ops::{Add, Mul, Sub, Div};
use num_traits::{Zero, One};

pub trait PrimeField:
    Zero + One + Add + Mul + ModReduce + Sub + Div + Eq {}

pub trait ModReduce {
    #[must_use]
    fn mod_reduce(product: u128) -> u64;
}
```

# Inefficient prime-field arithmetic in Rust

```
// 2^64 - 2^32 + 1
pub const P64: u64 = 0xffff_ffff_0000_0001;
pub const P128: u128 = 0xffff_ffff_0000_0001;

pub fn add(x: u64, y: u64) -> u64 {
    let sum: u128 = x as u128 + y as u128;
    (sum % P128) as u64
}

pub fn mul(x: u64, y: u64) -> u64 {
    let product: u128 = x as u128 * y as u128;
    (product % P128) as u64
}
```

# Efficient prime-field +

```
pub fn add_fast(x: u64, y: u64) -> u64 {  
    let mut sum: u128 = x as u128 + y as u128;  
    if sum > P128 {  
        sum -= P128;  
    }  
    sum as u64  
}
```

# Efficient prime-field $\times$

$$(for\ p = 2^{64} - 2^{32} + 1)$$

Credit: [cp4space.hatsya.com](https://cp4space.hatsya.com)'s blog post:

An efficient prime for number-theoretic transforms

- Elements in  $\mathbb{F}_p$  where  $p = 2^{64} - 2^{32} + 1$  fit nicely inside a 64-bit machine word, and "mod  $p$ " is possible without multiplication or division.

# Efficient prime-field $\times$

$$(for\ p = 2^{64} - 2^{32} + 1)$$

Any non-negative integer less than  $2^{159}$  can be written as  $A2^{96} + B2^{64} + C$  where A is a 63-bit integer, B is a 32-bit integer, and C is a 64-bit integer.

# Efficient prime-field $\times$

$$(for\ p = 2^{64} - 2^{32} + 1)$$

Since  $2^{96}$  is congruent to  $-1$  modulo  $p$ , this can be rewritten as  $B2^{64} + (C - A)$ . If  $A > C$ ,  $B2^{64}$  could underflow, in which case we can add  $p$ , resulting in a 96-bit integer.

# Efficient prime-field $\times$

$$(for\ p = 2^{64} - 2^{32} + 1)$$

To reduce this to a 64-bit integer,  $2^{64}$  is congruent to  $2^{32} - 1$ , so we can multiply  $B$  by  $2^{32}$  using a binary shift and a subtraction, and then add it to the result. We might encounter an overflow, but we can correct for that by subtracting  $p$ .

# Efficient prime-field $\times$

$$(for\ p = 2^{64} - 2^{32} + 1)$$

See the [source code](#). How fast is it, then?



# cargo criterion

11th Gen Intel(R) Core(TM) i7-11700 @ 2.50GHz

add/baseline/1000	time:	[0.0000 ps 0.0000 ps 0.0000 ps]
add/mod/1000	time:	[2.5268 μs 2.5290 μs 2.5317 μs]
add/fast/1000	time:	[1.0189 μs 1.0195 μs 1.0202 μs]
add/winterfell/1000	time:	[1.0564 μs 1.0573 μs 1.0581 μs]
mul/baseline/1000	time:	[0.0000 ps 0.0000 ps 0.0000 ps]
mul/mod/1000	time:	[2.5743 μs 2.5756 μs 2.5770 μs]
mul/reduce159/1000	time:	[1.6030 μs 1.6041 μs 1.6052 μs]
mul/reduce_montgomery/1000	time:	[1.3197 μs 1.3206 μs 1.3215 μs]

# cargo criterion

MacBook Pro M1 2021

add/mod/1000	time:	[7.2183 μs 7.2243 μs 7.2303 μs]
add/fast/1000	time:	[1.3538 μs 1.3570 μs 1.3602 μs]
add/winterfell/1000	time:	[1.2872 μs 1.2876 μs 1.2880 μs]
mul/mod/1000	time:	[13.447 μs 13.467 μs 13.487 μs]
mul/reduce159/1000	time:	[1.0562 μs 1.0569 μs 1.0576 μs]
mul/reduce_montgomery/1000	time:	[975.32 ns 975.60 ns 975.94 ns]

# Godbolt!

See the disassembly for [x86\\_64](#) and [aarch64](#).

# That was $+$ and $\times$ ... what else?

For prime fields and uni-/multivariate polynomials:

- `AddAssign` ( `+=` ), `MulAssign` ( `*=` ), when possible
- `.mod_pow()` with halving strategy ( $x^{2n} = x^n \cdot x^n$ )
- Generally, don't divide, multiply by inverses.
- Specialize `.square()`.
- Only re-allocate coefficients when `rhs.len() > lhs.len()`.
- Short-circuit polynomial multiplication when operand is  $f(x) = 0$  or  $f(x) = 1$ .

# EOF

See Alan Szepieniec's STARK anatomy tutorial:  
<https://neptune.cash/learn/stark-anatomy/>